

DD2424 Assignment 3



Wei Wang

wewang@kth.se

Royal Institute of Technology

DD2424 Deep Learning in Data Science

Assignment Report

Introduction

The object of this assignment is to upgrade the 2-layer networks into k-layer networks with multiple outputs to classify images (once again) from the CIFAR-10 dataset. In order to get better performance, batch normalization has to be incorporated into the k-layer network both for training and testing. After adding batch normalization, parameter searching is also required for getting better outcome.

Exercise

Exercise 1

Upgrade Assignment 2 code to train & test k-layer networks

In order to successfully test k-layer networks. Here are several subfunctions I have to write:

1. ***[W,b,GDparams]= ParameterInit3(train_X,...,alpha)***

Initialize almost all necessary parameter and generate W and b according to "layers" and "hidden_nodes"

2. ***EvaluateClassifier3(X,W,b,GDparams,varargin)***

Take the input data X and make the forward pass processing according to the W b. The intermediary score h, S \hat{S} and h are recorded. Then [P,S,S_hat,h,M_bn,V_bn] are returned. (M_bn,V_bn are used for batch normalization)

3. ***ComputeGradients3(X,...,V_bn)***

Compute the gradient of the cost function for using the gradient equations in the lectures notes 4. The snippet of code of computegradient3 is shown as Listing 1. Then I wrote a short script to compute the differences between analytic and the numerical gradients. As shown in Figure 2. I set the batch to be 20 ($Ttrain_X = train_X(:, 1 : 20); Ttrain_Y = train_Y(:, 1 : 20)$) and δ in ComputeGradsNumSlow to be e^{-6} . A 3-layers nets is used for analytic analysis, the hidden nodes of layer-2 is 50 and the hidden nodes of layer-3 is 30. The outcome of diff is shown in Figure 1. we can see that the larger of the layers has less difference with numerical gradients. The former the layer is, the larger the difference seems to be. I also calculated the 2-layers nets gradients. The batch size of data is 20 and the hidden nodes of layer-2 is 50. The outcome is shown in Figure 3. We can also notice that when nets get deeper, the analysis gradient is not so precise.

decay_rate	1	
diff_b1	1.7351e-04	44
diff_b2	0.0115	45
diff_b3	2.1502e-09	46
diff_W1	1.1407e-04	47
diff_W2	2.6690e-05	48
diff_W3	3.3627e-05	49

Figure 1: the diff of 3-layer nets without batch normalization

decay_rate	1
diff_b1	6.3824e-07
diff_b2	1.0663e-09
diff_W1	7.7567e-07
diff_W2	1.0936e-07

Figure 2: the diff of 2-layer nets without batch normalization

```

1  for i = 1:N
2      Yn = Y(:, i); % (K x 1)
3      Pn = P(:, i); % (K x 1)
4      Xn = X(:, i); % (d x 1)
5      % g = - Yn' / (Yn'*Pn) * (diag(Pn) - Pn*Pn');
6      g = -(Yn-Pn)';
7      % gradient L w.r.t b{L} = g
8      for j = (L):-1:2
9          grad_b{j} = grad_b{j} + g';
10         grad_W{j} = grad_W{j} + g'*h{j-1}(:, i)';
11         % update g
12         % (1 x m)
13         g = g*W{j};
14         g = g*diag(S{j-1}(:, i) > 0);
15     end
16     % (1 x m) —> (1 x m x m x m)
17     grad_b{1} = grad_b{1} + g';
18     grad_W{1} = grad_W{1} + g'*Xn';
19 end
20 % gradient J
21 for i = 1:L
22     grad_W{i} = (1/N)*grad_W{i} + 2*lambda*W{i};
23     grad_b{i} = (1/N)*grad_b{i};
24 end
25 end

```

Listing 1: snippet of ComputeGradients3

4. *MiniBatchGD3(X, Y, ...)* and *Run3(trainP, validP, GDparams, W, b)*

MiniBatch3 and Run3 are upgraded from assignment 2. GDparams carries

many necessary parameters such as GDparams.IFbn to decide whether use batch normalization or not.

Exercise 2

Can I train a 3-layer network?

2-layers without batch normalization:

Data_batch.1.mat is used for training and test_batch.mat is used for test. Using the parameters in assignment2 by running 20 epochs. The test accuracy is around 0.44 after 20 epochs showing that the upgrading succeed.

```

1 ##### epoch= 17 #####
2 train loss=1.429881 Validation loss = 1.671364
3 train acc=0.530700 Validation acc = 0.437600
4 ##### epoch= 18 #####
5 train loss=1.419637 Validation loss = 1.670553
6 train acc=0.533800 Validation acc = 0.438700
7 ##### epoch= 19 #####
8 train loss=1.410002 Validation loss = 1.669741
9 train acc=0.537300 Validation acc = 0.441700
10 ##### epoch= 20 #####
11 train loss=1.401479 Validation loss = 1.669319
12 train acc=0.541100 Validation acc = 0.444300

```

Listing 2: 2-layer network without bn running outcome

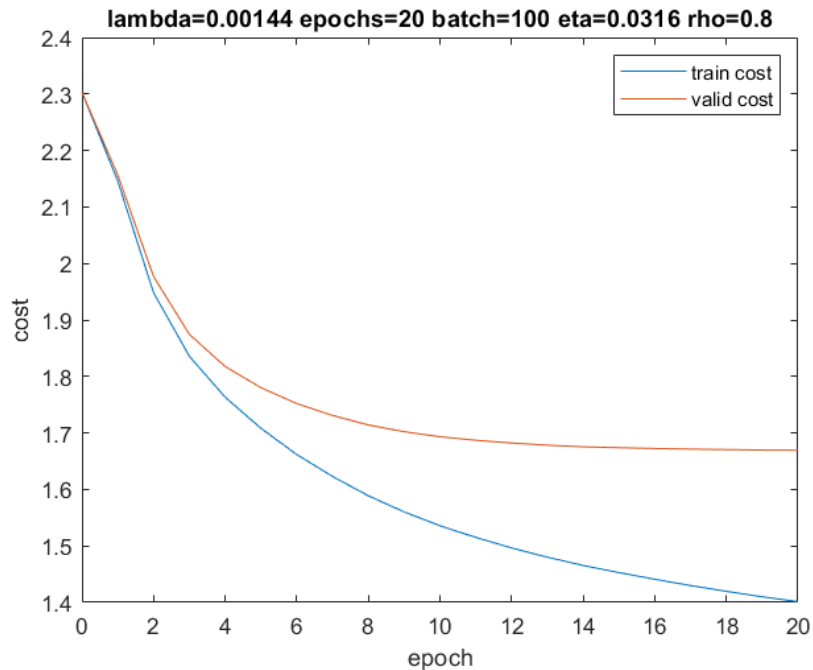


Figure 3: training curve of 2-layers network without bn

3-layers without batch normalization:

When adding 1 more layer to the network, the network seems stop learning using the same parameters as the previous one. When eta is increased, training becomes

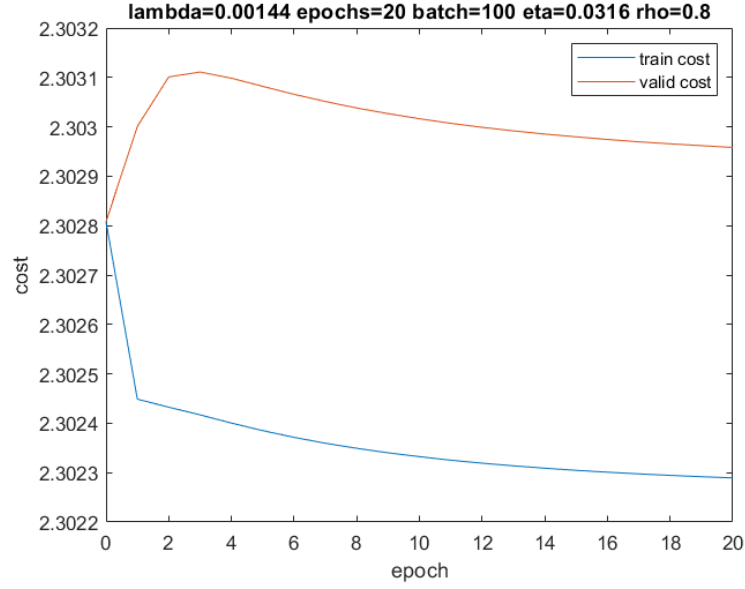


Figure 4: training curve of 3-layers network without bn

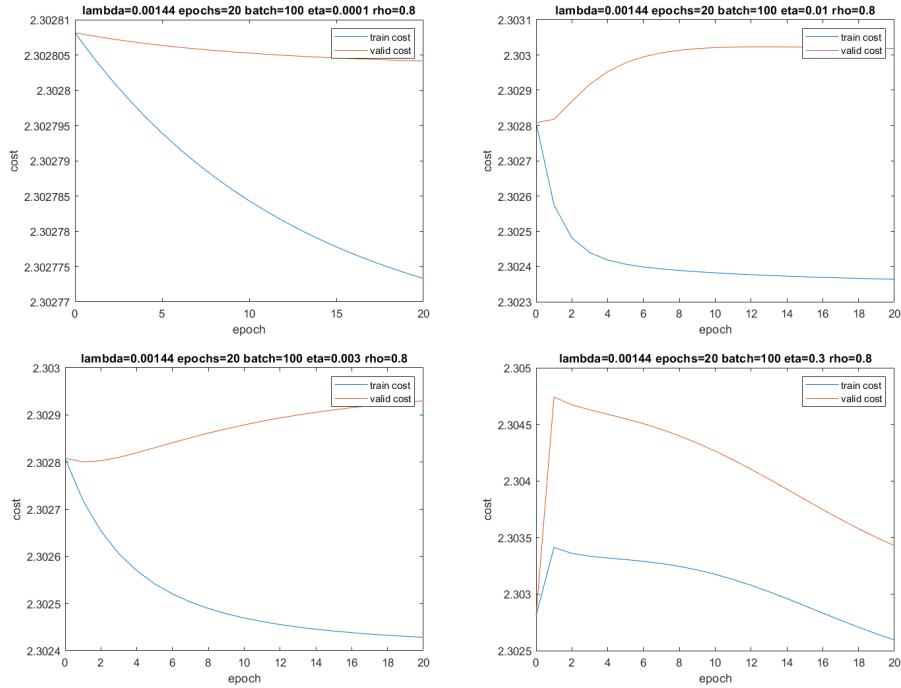


Figure 5: training curve of 3-layers network without bn

more unstable. When eta is decreased, the curve became more stable but the loss can

hardly get smaller. I got the conclusion that **when using bigger eta, the network which is without BN will be more unstable and hard for training.** After parameter searching, a better outcome is shown in Figure7. we can notice the curve is not stable.

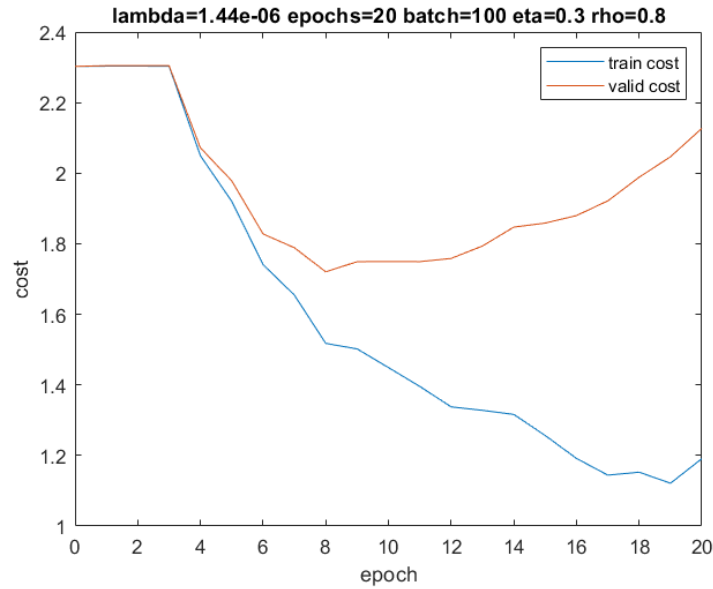


Figure 6: A better training curve of 3-layers network without bn

```

1 ##### epoch= 17 #####
2 train loss=1.144251 Validation loss = 1.921561
3 train acc=0.605100 Validation acc = 0.399400
4 ##### epoch= 18 #####
5 train loss=1.152620 Validation loss = 1.988506
6 train acc=0.597000 Validation acc = 0.391600
7 ##### epoch= 19 #####
8 train loss=1.121293 Validation loss = 2.046439
9 train acc=0.610200 Validation acc = 0.390900
10 ##### epoch= 20 #####
11 train loss=1.190905 Validation loss = 2.126961
12 train acc=0.594200 Validation acc = 0.379600

```

Listing 3: 3-layer network without bn running outcome

Exercise 3

1. $[s_bn, mean_scores, var_scores] = \text{BatchNormalization}(scores, varargin)$

I wrote the Batchnormalization() forward pass function according to equations (11) -(17) in Lab3 instruction. Is is assumed that the layer means and variances are computed from the mini-batch data sent into the function while training,

and the un-normalized scores are normalized by known pre-computed means and variances while testing. I use the **varargin** cell structure to control this. The Batchnormalization() is shown below:

```

1  function [s_bn, mean_scores, var_scores] = BatchNormalization(
    scores, varargin)
2  if numel(varargin) == 2
3      eps = 1e-6;
4      mean_scores = varargin{1};
5      var_scores = varargin{2};
6      s_bn = diag(var_scores+eps)^(-0.5) *(scores-repmat(
    mean_scores,1,size(scores,2)));
7  else
8      eps = 1e-6;
9      n = size(scores,2);
10     mean_scores = mean(scores,2);
11     var_scores = var(scores, 0, 2);
12     var_scores = var_scores *(n-1)/n;
13     s_bn = diag(var_scores+eps)^(-0.5) *(scores-repmat(
    mean_scores,1,size(scores,2)));
14 end
15 end

```

Listing 4: snippet Batchnormalization

2. *EvaluateClassifier3(X,W,b,GDparams,varargin)*

The Evaluateclassifier3() is also changed according to Batchnormalization():

```

1  for i = 1:L-1
2      if IFbn
3          if numel(varargin) == 2
4              mean_score = varargin{1}{i};
5              var_score = varargin{2}{i};
6              S_hat{i} = S{i};
7              [S{i},M_bn{i},V_bn{i}] = BatchNormalization(S{i},
    mean_score, var_score);
8          else
9              S_hat{i} = S{i};
10             [S{i},M_bn{i},V_bn{i}] = BatchNormalization(S{i});
11         end
12     end
13     h{i} = max(0,S{i});
14     S{i+1} = W{i+1}*h{i}+repmat(b{i+1},1,n);
15 end

```

Listing 5: snippet Batchnormalization

3. *ComputeGradients3(X,h,S,S_hat,Y, P, W,...,V_bn)*

I wrote the computeGradient3() into two parts: one is the pipeline without batch normalization, the other is the pipeline totally following the back-prop algorithm in Lab3

instruction. Then I used the parameter **Ifbn** to switch. The script is shown below:

```

1 if IFbn
2     g = cell(N,1);
3     % calculate gk
4     for i = 1:N
5         Yn = Y(:,i); % (K x 1)
6         Pn = P(:,i); % (K x 1)
7         % Xn = X(:,i); % (d x 1)
8         g{i} = - Yn' / (Yn'*Pn) * (diag(Pn) - Pn*Pn');
9         % g{i} = -(Yn-Pn)';
10        % gradient L w.r.t b{L} = g
11        grad_b{L} = grad_b{L} + g{i}';
12        grad_W{L} = grad_W{L} + g{i}'*h{L-1}(:,i)';
13    end
14    % get grad_bk grad_wk
15    grad_b{L} = grad_b{L}/N;
16    grad_W{L} = grad_W{L}/N+2*lambda*W{L};
17    % propagate to previous layers
18    for i = 1:N
19        g{i} = g{i}*W{L};
20        g{i} = g{i}*diag(S{L-1}(:,i)>0);
21    end
22    % bn
23    for i = L-1:-1:2
24        g = BatchNormBackPass(g, S_hat{i}, M_bn{i}, V_bn{i});
25        for j = 1:N
26            grad_b{i} = grad_b{i} + g{j}';
27            grad_W{i} = grad_W{i} + g{j}'*h{i-1}(:,j)';
28        end
29        grad_b{i} = grad_b{i}/N;
30        grad_W{i} = grad_W{i}/N+2*lambda*W{i};
31
32        for m = 1:N
33            g{m} = g{m}*W{i};
34            g{m} = g{m}*diag(S{i-1}(:,m)>0);
35        end
36    end
37    g = BatchNormBackPass(g, S_hat{1}, M_bn{1}, V_bn{1});
38    for j = 1:N
39        grad_b{1} = grad_b{1} + g{j}';
40        grad_W{1} = grad_W{1} + g{j}'*X(:,j)';
41    end
42    grad_b{1} = grad_b{1}/N;
43    grad_W{1} = grad_W{1}/N+2*lambda*W{1};
44 else
45     ... back-prop without BN
46
47 end

```

Listing 6: snippet of ComputeGradients3

4. $g = \text{BatchNormBackPass}(g, S, \mu, \text{var})$

I wrote the BatchNormBackpass following the equation in the last page of Lecture note 4.

5. *MiniBatchGD3(X, Y, GDparams, W, b, varargin)*

In MiniBatch(), I add exponential moving average for batch means and variances.

```

1  for j=1:fix(N/n_batch)
2      j_start = (j-1)*n_batch + 1;
3      j_end = j*n_batch;
4      inds = j_start:j_end;
5      Xbatch = Xtrain(:, inds);
6      Ybatch = Ytrain(:, inds);
7      [P,S,S_hat,h,M_bn,V_bn] = EvaluateClassifier3(Xbatch,Wstar,
      bstar,GDparams);
8      if GDparams.IFbn
9          if numel(varargin) == 0 && j == 1
10             M_av = M_bn;
11             V_av = V_bn;
12         else
13             for i = 1:L-1
14                 M_av{i} = GDparams.alpha*M_av{i} +(1-GDparams.
alpha)*M_bn{i};
15                 V_av{i} = GDparams.alpha*V_av{i} +(1-GDparams.
alpha)*V_bn{i};
16             end
17         end
18
19         [grad_W,grad_b] = ComputeGradients3(Xbatch,h,S,S_hat,
Ybatch, P, Wstar, GDparams,M_bn,V_bn);
20     else
21         [grad_W,grad_b] = ComputeGradients3(Xbatch,h,S,S_hat,
Ybatch, P, Wstar, GDparams);
22     end

```

Listing 7: snippet MiniBatchGD3

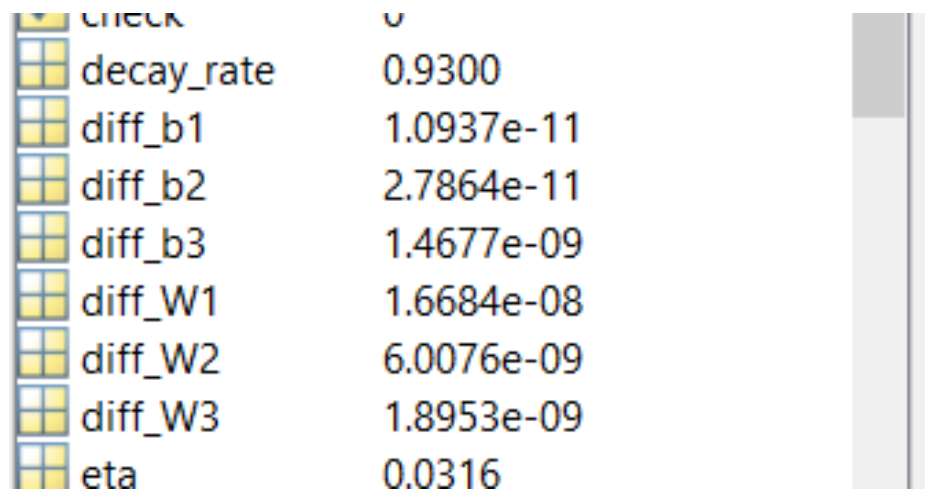
The M_av and V_av were initialized to be M_bn and V_bn at the very first epoch.

Then I computed the gradient both in numerical calculation and analytic calculation. The batch of input data(Ttrain_X = train_X(:,1:20)) is 20. In the outcome. we can see the difference of W3 is the smallest.

Task

1. State how you checked your analytic gradient computations and whether you think that your gradient computations are bug free for your k-layer network with batch normalization.

The pipeline of back-prop of BN is based on the instruction equation 19-26. In order to make sure there is bug free, I calculate the gradient difference between numerical calculation and analytic calculation. The difference showing that my gradient computations are bug free.



decay_rate	0.9300
diff_b1	1.0937e-11
diff_b2	2.7864e-11
diff_b3	1.4677e-09
diff_W1	1.6684e-08
diff_W2	6.0076e-09
diff_W3	1.8953e-09
eta	0.0316

Figure 7: The gradient differences of 3-layers network with BN

2. Include graphs of the evolution of the loss function when you tried to train your 3-layer network without batch normalization and with batch normalization.

1. The parameters of 3-layers without BN:

when I use the parameters searched in assignment2, the 3-layers without BN can hardly converge.

rng_number=40; n_epochs=20; eta=0.0316; lambda=1.46e-4; rho=0.88;
decay_rate=0.93

The training curve is shown in Figure 8.

```
##### epoch= 17 #####
1 train loss=2.302298 Validation loss = 2.302970
2 train acc=0.103200 Validation acc = 0.101000
3 ##### epoch= 18 #####
4 train loss=2.302295 Validation loss = 2.302966
5 train acc=0.103200 Validation acc = 0.101000
6
```




Figure 8: The training curve of 3-layers network without BN

```

7 ##### epoch= 19 #####
8 train loss=2.302292 Validation loss = 2.302962
9 train acc=0.103200 Validation acc = 0.101000
10 ##### epoch= 20 #####
11 train loss=2.302290 Validation loss = 2.302958
12 train acc=0.103200 Validation acc = 0.101000

```

Listing 8: training loss without BN

After a rough search, the training curve of 3-layer network without BN is shown below Figure 9. The loss is shown in Listing 9. I changed the parameters to be :

rng_number=40; n_epochs=20; eta=0.0316; lambda=1.46e-6; rho=0.8; decay_rate=0.93

```

1 ##### epoch= 17 #####
2 train loss=1.144251 Validation loss = 1.921561
3 train acc=0.605100 Validation acc = 0.399400
4 ##### epoch= 18 #####
5 train loss=1.152620 Validation loss = 1.988506
6 train acc=0.597000 Validation acc = 0.391600
7 ##### epoch= 19 #####
8 train loss=1.121293 Validation loss = 2.046439
9 train acc=0.610200 Validation acc = 0.390900
10 ##### epoch= 20 #####
11 train loss=1.190905 Validation loss = 2.126961
12 train acc=0.594200 Validation acc = 0.379600

```

Listing 9: training loss without BN




Figure 9: The training curve of 3-layers network without BN

2. The parameters of 3-layers with BN:

The parameter is the same as assignment2. It is clearly shown that using BN **can accelerate the convergence speed and make the training more stable**. Also, the outcome is better than the network without BN.

rng_number=40; n_epochs=20; eta=0.0316; lambda=1.46e-4; rho=0.88; decay_rate=0.93

```

1 ##### epoch= 17 #####
2 train loss=1.167175 Validation loss = 1.977223
3 train acc=0.647600 Validation acc = 0.416300
4 ##### epoch= 18 #####
5 train loss=1.179972 Validation loss = 2.004503
6 train acc=0.641900 Validation acc = 0.409200
7 ##### epoch= 19 #####
8 train loss=1.187421 Validation loss = 2.046314
9 train acc=0.643500 Validation acc = 0.407200
10 ##### epoch= 20 #####
11 train loss=1.134107 Validation loss = 2.053281
12 train acc=0.665400 Validation acc = 0.412800

```

Listing 10: training loss with BN

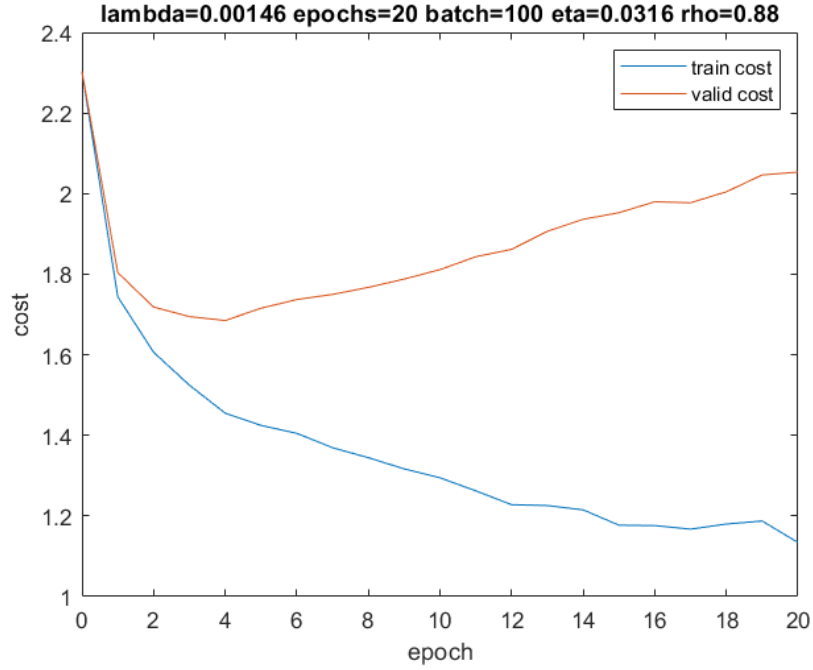


Figure 10: The gradient differences of 3-layers network with BN

3.State the range of the values you searched for lambda and eta, the number of epochs used for training during the ne search, and the hyper-parameter settings for your best performing 3-layer network you trained with batch normalization. Also state the test accuracy achieved by network.

First rough search:

I set the *search number* = 100; *rng_number*=40; *n_epochs*=5; $\eta = e^{-4} - e^{-1}$; $\lambda = e^{-7} - e^{-1}$; $\rho = 0.9$; *decay_rate*=0.95; *train data* = *data_batch_1.mat*; *test data* = *test_batch.mat*

The top 3 graph is shown below




Figure 11: The Top-3 learning curve (rough search)

Table 1: Top16 coarse searching testing accuracy with 5 epochs

	eta	lambda	train accuracy	test accuracy
top1	0.021909942	2.81E-05	0.531	0.4272
top2	0.020810752	8.80E-07	0.5293	0.4257
top3	0.026202212	8.92E-07	0.5381	0.4249
top4	0.01653527	4.13E-06	0.5239	0.4237
top5	0.02430095	0.001323581	0.5321	0.4217
top6	0.103422378	0.000444255	0.5321	0.421
top7	0.007062768	7.66E-07	0.5137	0.4207
top8	0.01722279	9.63E-05	0.529	0.4204
top9	0.007056148	0.098428852	0.5121	0.4198
top10	0.055774277	0.024148088	0.5283	0.4195
top11	0.072874748	0.002220315	0.5396	0.4191
top12	0.043335998	0.000159417	0.53	0.4191
top13	0.045183708	0.000234025	0.5403	0.4189
top14	0.034854939	0.004073542	0.5193	0.4188
top15	0.209136611	0.000114808	0.5183	0.4187
top16	0.02837389	3.63E-06	0.52	0.418

Precise search:

I set the *search number* = 100; *rng_number*=40; *n_epochs*=10; *eta*= 0.009 – 0.035; *lambda*= $e^{-7} - e^{-3}$; *rho*=0.9; *decay_rate*=0.95; *train data*=*data_batch_1.mat*; *test data* = *test_batch.mat*

The Best test accuracy at the second search is 0.435 (10 epoch)

Table 2: Top16 precise searching testing accuracy with 5 epochs

	eta	lambda	train accuracy	test accuracy
top1	0.009976828	2.24E-06	0.5944	0.4351
top2	0.015179735	1.76E-05	0.5934	0.4285
top3	0.022828666	1.68E-06	0.6111	0.428
top4	0.020718543	4.91E-05	0.5971	0.4277
top5	0.029637711	7.51E-07	0.6012	0.4268
top6	0.031951095	3.66E-05	0.6041	0.4267
top7	0.01832903	1.16E-06	0.5943	0.4267
top8	0.013022228	4.47E-07	0.5985	0.4264
top9	0.013763121	2.90E-06	0.5978	0.426
top10	0.022244822	6.03E-07	0.5875	0.4259
top11	0.017005871	1.15E-07	0.5895	0.4257
top12	0.021288386	3.99E-06	0.5997	0.4255
top13	0.017281191	1.59E-05	0.5861	0.4251
top14	0.027226389	9.58E-07	0.6132	0.4246
top15	0.016494211	7.45E-05	0.5926	0.4243
top16	0.034853878	6.75E-06	0.603	0.4235




Figure 12: The Top-3 learning curve (precise search)

4. Plot the training and validation loss for your 2-layer network with batch normalization with 3 different learning rates (small, medium, high) for 10 epochs and make the same plots for a 2-layer network with no batch normalization.

After a rough search (*search number = 20*), I set the parameters to be:

rng_number=40; n_epochs=10; lambda=1.2617e⁻⁴; rho=0.9; decay_rate=0.95; train data=data_batch_1.mat; test data = test_batch.mat

1. *eta = 0.3(high)*




Figure 13: (high lr)Left:BN Right:without BN

```

1 % with BN
2 ##### epoch= 10 #####
3 train loss=1.304410 Validation loss = 2.027296
4 train acc=0.589400 Validation acc = 0.395000
5 % without BN
6 bad parameter, too large
7 Elapsed time is 6.014476 seconds.

```

Listing 11: training loss

2. *eta = 0.03(medium)*




Figure 14: (medium lr),Left:BN Right:without BN

```

1 % with BN
2 ##### epoch= 10 #####
3 train loss=1.201245 Validation loss = 1.699140
4 train acc=0.595900 Validation acc = 0.421600
5 % without BN
6 ##### epoch= 10 #####
7 train loss=1.265019 Validation loss = 1.639208
8 train acc=0.567800 Validation acc = 0.432900

```

Listing 12: training loss

3. $\eta = 0.003$ (small)




Figure 15: (small lr),Left:BN Right:without BN

```

1 % with BN
2 ##### epoch= 10 #####
3 train loss=1.455016 Validation loss = 1.709162
4 train acc=0.514200 Validation acc = 0.399000
5 % without BN
6 ##### epoch= 10 #####
7 train loss=1.985748 Validation loss = 2.013167

```



```
8 train acc=0.262800 Validation acc = 0.246800
```

Listing 13: training loss

Conclusion

From these comparison above, I can draw the conclusion that incorporating Batch Normalization can make the network training more stable and converge faster. It can adapt larger range of eta without collapse.